

Problem Set 3: Game of Fifteen

This is CS50. Harvard University. Fall 2015.

Table of Contents

[Objectives](#)

[Recommended Reading](#)

[Academic Honesty](#)

[Reasonable](#)

[Not Reasonable](#)

[Getting Ready](#)

[Getting Started](#)

[Searching](#)

[search](#)

[Sorting](#)

[sort](#)

[search](#)

[The Game Begins](#)

[questions](#)

[fifteen](#)

[How to Submit](#)

[Step 1 of 2](#)

[Step 2 of 2](#)

Objectives

- Accustom you to reading someone else's code.
- Introduce you to larger programs and programs with multiple source files.
- Empower you with Makefiles.
- Implement a party favor.

Recommended Reading

- Page 17 of <http://www.howstuffworks.com/c.htm> (<http://www.howstuffworks.com/c.htm>).
- Chapters 20 and 23 of *Absolute Beginner's Guide to C*.
- Chapters 13, 15, and 18 of *Programming in C*.

Academic Honesty

This course's philosophy on academic honesty is best stated as "be reasonable." The course recognizes that interactions with classmates and others can facilitate mastery of the course's material. However, there remains a line between enlisting the help of another and submitting the work of another. This policy characterizes both sides of that line.

The essence of all work that you submit to this course must be your own. Collaboration on problem sets is not permitted except to the extent that you may ask classmates and others for help so long as that help does not reduce to another doing your work for you. Generally speaking, when asking for help, you may show your code to others, but you may not view theirs, so long as you and they respect this policy's other constraints. Collaboration on quizzes is not permitted at all. Collaboration on the course's final project is permitted to the extent prescribed by its specification.

Below are rules of thumb that (inexhaustively) characterize acts that the course considers reasonable and not reasonable. If in doubt as to whether some act is reasonable, do not commit it until you solicit and receive approval in writing from the course's heads. Acts considered not reasonable by the course are handled harshly. If the course refers some matter for disciplinary action and the outcome is punitive, the course reserves the right to impose local sanctions on top of that outcome that may include an unsatisfactory or failing grade for work submitted or for the course itself.

If you commit some act that is not reasonable but bring it to the attention of the course's heads within 72 hours, the course may impose local sanctions that may include an unsatisfactory or failing grade for work submitted, but the course will not refer the matter for further disciplinary action except in cases of repeated acts.

Reasonable

- Communicating with classmates about problem sets' problems in English (or some other spoken language).
- Discussing the course's material with others in order to understand it better.
- Helping a classmate identify a bug in his or her code at office hours, elsewhere, or even online, as by viewing, compiling, or running his or her code, even on your own computer.
- Incorporating snippets of code that you find online or elsewhere into your own code, provided that those snippets are not themselves solutions to assigned problems and that you cite the snippets' origins.
- Reviewing past semesters' quizzes and solutions thereto.
- Sending or showing code that you've written to someone, possibly a classmate, so that he or she might help you identify and fix a bug.
- Sharing snippets of your own code online so that others might help you identify and fix a bug.
- Turning to the web or elsewhere for instruction beyond the course's own, for references, and for solutions to technical difficulties, but not for outright solutions to problem set's problems or your own final project.
- Whiteboarding solutions to problem sets with others using diagrams or pseudocode but not actual code.
- Working with (and even paying) a tutor to help you with the course, provided the tutor does not do your work for you.

Not Reasonable

- Accessing a solution to some problem prior to (re-)submitting your own.
- Asking a classmate to see his or her solution to a problem set's problem before (re-)submitting your own.
- Decompiling, deobfuscating, or disassembling the staff's solutions to problem sets.
- Failing to cite (as with comments) the origins of code or techniques that you discover outside of the course's own lessons and integrate into your own work, even while respecting this policy's other constraints.
- Giving or showing to a classmate a solution to a problem set's problem when it is he or she, and not you, who is struggling to solve it.
- Looking at another individual's work during a quiz.

- Paying or offering to pay an individual for work that you may submit as (part of) your own.
 - Providing or making available solutions to problem sets to individuals who might take this course in the future.
 - Searching for, soliciting, or viewing a quiz's questions or answers prior to taking the quiz.
 - Searching for or soliciting outright solutions to problem sets online or elsewhere.
 - Splitting a problem set's workload with another individual and combining your work.
 - Submitting (after possibly modifying) the work of another individual beyond allowed snippets.
 - Submitting the same or similar work to this course that you have submitted or will submit to another.
 - Submitting work to this course that you intend to use outside of the course (e.g., for a job) without prior approval from the course's heads.
 - Using resources during a quiz beyond those explicitly allowed in the quiz's instructions.
 - Viewing another's solution to a problem set's problem and basing your own solution on it.
-

Getting Ready

First, re-acquaint yourself with with our old friends, linear search and binary search (and Patrick), if hazy:

 1/2 Linear Search



Next, get to know our new friends, bubble sort, selection sort, insertion sort, and merge sort (and Jackson, Tommy, and Rob):



Phew, so many shorts. And so many sorts! Ha.

Finally, a more in-depth look at debugging techniques from Dan! (Odds are these 23 minutes with Dan will save you hours over the course of the term, since GDB is a far better tool than `printf` in many cases!)

Section: Debugging with Dan Armendariz



Be sure you're reasonably comfortable answering the below when it comes time to submit this problem set's form!

- Why does binary search require that an array be sorted?
- Why is bubble sort in $O(n^2)$?
- Why is insertion sort in $\Omega(n)$?
- In no more than 3 sentences, how does selection sort work?
- What's an upper bound on the (worst-case) running time of merge sort?
- GDB lets you "debug" program, but, more specifically, what does it let you do?

Getting Started

Recall that, for Problem Sets 1 and 2, you started writing programs from scratch, creating your own `pset1` and `pset2` directories with `mkdir`. (And for Problem Set 0, you started writing programs in Scratch! Ha.) For Problem Set 3, you'll instead download "distribution code" (otherwise known as a "distro"), written by us, and add your own lines of code to it. You'll first need to read and understand our code, though, so this problem set is as much about learning to read someone else's code as it is about writing your own!

Let's get you started. Log into cs50.io (<https://cs50.io/>) and execute

```
update50
```

within a terminal window to make sure your workspace is up-to-date. If you somehow closed your terminal window (and can't find it!), make sure that **Console** is checked under the **View** menu, then click the green, circled plus (+) in CS50 IDE's bottom half, then select **New Terminal**.

Next, execute

```
cd ~/workspace
```

at your prompt to ensure that you're inside of `workspace` (which is inside of your home directory). Then execute

```
wget http://cdn.cs50.net/2015/fall/psets/3/pset3/pset3.zip
```

to download a ZIP of this problem set's distro into your workspace (with a command-line program called `wget`). You should see a bunch of output followed by:

```
'pset3.zip' saved
```

Confirm that you've indeed downloaded `pset3.zip` by executing

```
ls
```

and then run

```
unzip pset3.zip
```

to unzip the file. If you then run `ls` again, you should see that you have a newly unzipped directory called `pset3` as well. Proceed to execute

```
cd pset3
```

followed by

```
ls
```

and you should see that the directory contains two "subdirectories":

```
fifteen find
```

How fun!

Searching

Okay, let's dive into one of those subdirectories. Execute the command below in a terminal window.

```
cd ~/workspace/pset3/find
```

If you list the contents of this directory (remember how?), you should see the below.

```
helpers.c helpers.h Makefile find.c generate.c
```

Wow, that's a lot of files, eh? Not to worry, we'll walk you through them.

Implemented in `generate.c` is a program that uses a "pseudorandom-number generator" (via a function called `drand48`) to generate a whole bunch of random (well, pseudorandom, since computers can't actually generate truly random) numbers, one per line. Go ahead and compile this program by executing the command below.

```
make generate
```

Now run the program you just compiled by executing the command below.

```
./generate
```

You should be informed of the program's proper usage, per the below.

```
Usage: generate n [s]
```

As this output suggests, this program expects one or two command-line arguments. The first, `n`, is required; it indicates how many pseudorandom numbers you'd like to generate. The second, `s`, is optional, as the brackets are meant to imply; if supplied, it represents the value that the pseudorandom-number generator should use as its "seed." A seed is simply an input to a pseudorandom-number generator that influences its outputs. For instance, if you seed `drand48` by first calling `srand48` (another function whose purpose is to "seed" `drand48`) with an argument of, say, `1`, and then call `drand48` itself three times, `drand48` might

return 2728, then 29785, then 54710. But if you instead seed drand48 by first calling srand48 with an argument of, say, 2, and then call drand48 itself three times, drand48 might instead return 59797, then 10425, then 37569. But if you re-seed drand48 by calling srand48 again with an argument of 1, the next three times you call drand48, you'll again get 2728, then 29785, then 54710! See, not so random.

Go ahead and run this program again, this time with a value of, say, 10 for n, as in the below; you should see a list of 10 pseudorandom numbers.

```
./generate 10
```

Run the program a third time using that same value for n; you should see a different list of 10 numbers. Now try running the program with a value for s too (e.g., 0), as in the below.

```
./generate 10 0
```

Now run that same command again:

```
./generate 10 0
```

Bet you saw the same "random" sequence of ten numbers again? Yup, that's what happens if you don't vary a pseudorandom number generator's initial seed.

Now take a look at generate.c itself. (Remember how?) Comments atop that file explain the program's overall functionality. But it looks like we forgot to comment the code itself. Read over the code carefully until you understand each line and then comment our code for us, replacing each TODO with a phrase that describes the purpose or functionality of the corresponding line(s) of code. (Know that an unsigned int is just an int that cannot be negative.) And for more details on rand and srand, recall that you can execute:

```
man drand48
```

and:

```
man srand48
```

Once done commenting generate.c, re-compile the program to be sure you didn't break anything by re-executing the command below.

```
make generate
```

If generate no longer compiles properly, take a moment to fix what you broke!

Now, recall that make automates compilation of your code so that you don't have to execute clang manually along with a whole bunch of switches. Notice, in fact, how make just executed a pretty long command for you, per the tool's output. However, as your programs grow in size, make won't be able to infer from context anymore how to compile your code; you'll need to start telling make how to compile your program, particularly when they involve multiple source (i.e., .c) files. And so we'll start relying on "Makefiles," configuration files that tell make exactly what to do.

How did `make` know how to compile `generate` in this case? It actually used a configuration file that we wrote. Go ahead and look at the file called `Makefile` that's in the same directory as `generate.c`. This `Makefile` is essentially a list of rules that we wrote for you that tells `make` how to build `generate` from `generate.c` for you. The relevant lines appear below.

```
generate: generate.c
    clang -ggdb3 -O0 -std=c11 -Wall -Werror -o generate generate.c
```

The first line tells `make` that the "target" called `generate` should be built by invoking the second line's command. Moreover, that first line tells `make` that `generate` is dependent on `generate.c`, the implication of which is that `make` will only re-build `generate` on subsequent runs if that file was modified since `make` last built `generate`. Neat time-saving trick, eh? In fact, go ahead and execute the command below again, assuming you haven't modified `generate.c`.

```
make generate
```

You should be informed that `generate` is already up-to-date. Incidentally, know that the leading whitespace on that second line is not a sequence of spaces but, rather, a tab. Unfortunately, `make` requires that commands be preceded by tabs, so be careful not to change them to spaces, else you may encounter strange errors! The `-Werror` flag, recall, tells `clang` to treat warnings (bad) as though they're errors (worse) so that you're forced (in a good, instructive way!) to fix them.

Now take a look at `find.c`. Notice that this program expects a single command-line argument: a "needle" to search for in a "haystack" of values. Once done looking over the code, go ahead and compile the program by executing the command below.

```
make find
```

Notice, per that command's output, that `make` actually executed the below for you.

```
clang -ggdb3 -O0 -std=c11 -Wall -Werror -o find find.c helpers.c -lcs50 -lm
```

Notice further that you just compiled a program comprising not one but two `.c` files: `helpers.c` and `find.c`. How did `make` know what to do? Well, again, open up `Makefile` to see the man behind the curtain. The relevant lines appear below.

```
find: find.c helpers.c helpers.h
    clang -ggdb3 -O0 -std=c11 -Wall -Werror -o find find.c helpers.c -lcs50 -lm
```

Per the dependencies implied above (after the colon), any changes to `find.c`, `helpers.c`, or `helpers.h` will compel `make` to rebuild `find` the next time it's invoked for this target.

Go ahead and run this program by executing, say, the below.

```
./find 13
```

You'll be prompted to provide some hay (i.e., some integers), one "straw" at a time. As soon as you tire of providing integers, hit ctrl-d to send the program an EOF (end-of-file) character. That character will compel `GetInt` from the CS50 Library to return `INT_MAX`, a constant that, per `find.c`, will compel `find` to stop prompting for hay. The program will then look for that needle in the hay you provided, ultimately reporting whether the former was found in the latter. In short, this program searches an array for some value. At least, it should, but it won't find anything yet! That's where you come in. More on your role in a bit.

It turns out you can automate this process of providing hay, though, by "piping" the output of `generate` into `find` as input. For instance, the command below passes 1,000 pseudorandom numbers to `find`, which then searches those values for `42`.


```
./generate 1000 | ./find 42
```

Note that, when piping output from `generate` into `find` in this manner, you won't actually see `generate`'s numbers, but you will see `find`'s prompts.

Alternatively, you can "redirect" `generate`'s output to a file with a command like the below.

```
./generate 1000 > numbers.txt
```

You can then redirect that file's contents as input to `find` with the command below.

```
./find 42 < numbers.txt
```

Let's finish looking at that `Makefile`. Notice the line below.

```
all: find generate
```

This target implies that you can build both `generate` and `find` simply by executing the below.

```
make all
```

Even better, the below is equivalent (because `make` builds a `Makefile`'s first target by default).

```
make
```

If only you could whittle this whole problem set down to a single command! Finally, notice these last lines in `Makefile`:

```
clean:
    rm -f *.o a.out core find generate
```

This target allows you to delete all files ending in `.o` or called `core` (more on that soon!), `find`, or `generate` simply by executing the command below.

```
make clean
```

Be careful not to add, say, `*.c` to that last line in `Makefile`! (Why?) Any line, incidentally, that begins with `#` is just a comment.

search

And now the fun begins! Notice that `find.c` calls `search`, a function declared in `helpers.h`. Unfortunately, we forgot to implement that function fully in `helpers.c`! (To be sure, we could have put the contents of `helpers.h` and `helpers.c` in `find.c` itself. But it's sometimes better to organize programs into multiple files, especially when some functions are essentially utility functions that might later prove useful to other programs as well, much like those in the CS50 Library.) Take a peek at `helpers.c` with, and you'll see that `search` always returns `false`, whether or not `value` is in `values`. Re-write `search` in such a way that it uses linear search, returning `true` if `value` is in `values` and `false` if `value` is not in `values`. Take care to return `false` right away if `n` isn't even positive.

When ready to check the correctness of your program, try running the command below.

```
./generate 1000 50 | ./find 127
```

Because one of the numbers outputted by `generate`, when seeded with `50`, is `127`, your code should find that "needle"! By contrast, try running the command below as well.

```
./generate 1000 50 | ./find 128
```

Because `128` is not among the numbers outputted by `generate`, when seeded with `50`, your code shouldn't find that needle. Best to try some other tests as well, as by running `generate` with some seed, taking a look at its output, then piping that same output to `find`, looking for a "needle" you know to be among the "hay".

Incidentally, note that `main` in `find.c` is written in such a way that `find` returns `0` if the needle is found, else it returns `1`. You can check the so-called "exit code" with which `main` returns by executing

```
echo $?
```

after running some other command. For instance, assuming your implementation of `search` is correct, if you run

```
./generate 1000 50 | ./find 127  
echo $?
```

you should see `0`, since `127` is, again, among the 1,000 numbers outputted by `generate` when seeded with `50`, and so `search` (written by you) should return `true`, in which case `main` (written by us) should return (i.e., exit with) `0`. By contrast, assuming your implementation of `search` is correct, if you run

```
./generate 1000 50 | ./find 128  
echo $?
```

you should see `1`, since `128` is, again, not among the 1,000 numbers outputted by `generate` when seeded with `50`, and so `search` (written by you) should return `false`, in which case `main` (written by us) should return (i.e., exit with) `1`. Make sense?

When ready to check the correctness of your program officially with `check50`, you may execute the below.

```
check50 2015.fall.pset3.find helpers.c
```

Incidentally, be sure not to get into the habit of testing your code with `check50` before testing it yourself. (And definitely don't get into an even worse habit of only testing your code with `check50`!) Suffice it to say `check50` doesn't exist in the real world, so running your code with your own sample inputs, comparing actual output against expected output, is the best habit to get into sooner rather than later. Truly, don't do yourself a long-term disservice!

Anyhow, if you'd like to play with the staff's own implementation of `find`, you may execute the below.

```
~cs50/pset3/find
```

Sorting

Alright, linear search is pretty meh. Recall from Week 0 that we can do better, but first we'd best sort that hay.

sort

Notice that `find.c` calls `sort`, a function declared in `helpers.h`. Unfortunately, we forgot to implement that function fully too in `helpers.c`! Take a peek at `helpers.c`, and you'll see that `sort` returns immediately, even though `find`'s `main` function does pass it an actual array.

Now, recall the syntax for declaring an array. Not only do you specify the array's type, you also specify its size between brackets, just as we do for `haystack` in `find.c`:

```
int haystack[MAX];
```

But when passing an array, you only specify its name, just as we do when passing `haystack` to `sort` in `find.c`:

```
sort(haystack, size);
```

(Why do you think we pass in the size of that array separately?)

When declaring a function that takes a one-dimensional array as an argument, though, you don't need to specify the array's size, just as we don't when declaring `sort` in `helpers.h` (and `helpers.c`):

```
void sort(int values[], int n);
```

Go ahead and implement `sort` so that the function actually sorts, from smallest to largest, the array of numbers that it's passed, in such a way that its running time is in $O(n^2)$, where n is the array's size. Odds are you'll want to implement bubble sort, selection sort, or insertion sort, if only because we discussed them in Week 3. Just realize that there's no one "right" way to implement any of those algorithms; variations abound. In fact, you're welcome to improve upon them as you see fit, so long as your implementation remains in $O(n^2)$. However, take care not to alter our declaration of `sort`. Its prototype must remain:

```
void sort(int values[], int n);
```

As this return type of `void` implies, this function must not return a sorted array; it must instead "destructively" sort the actual array that it's passed by moving around the values therein. As we'll discuss in Week 4, arrays are not passed "by value" but instead "by reference," which means that `sort` will not be passed a copy of an array but, rather, the original array itself.

Although you may not alter our declaration of `sort`, you're welcome to define your own function(s) in `helpers.c` that `sort` itself may then call.

We leave it to you to determine how best to test your implementation of `sort`. But don't forget that `printf` and GDB are your friends. And don't forget that you can generate the same sequence of pseudorandom numbers again and again by explicitly specifying `generate`'s seed. Before you ultimately submit, though, be sure to remove any such calls to `printf`, as we like our programs' outputs just they way they are!

Here's Zamyła with some tips:



And if you'd like to play with the staff's own implementation of `find`, you may execute the below.

```
~cs50/pset3/find
```

search


Now that `sort` (presumably) works, it's time to improve upon `search`, the other function that lives in `helpers.c`. Recall that your first version implemented linear search. Rip out the lines that you wrote earlier (sniff) and re-implement `search` as Binary Search, that divide-and-conquer strategy that we employed in Week 0. You are welcome to take an iterative approach (as with a loop) or a recursive approach (wherein a function calls itself). If you pursue the latter, though, know that you may not change our declaration of `search`, but you may write a new, recursive function (that perhaps takes different parameters) that `search` itself calls. When it comes time to submit this problem set, it suffices to submit this new-and-improved version of `search`; you needn't submit your original version that used linear search.

Here's Zamyra again:




The Game Begins

And now it's time to play. The Game of Fifteen is a puzzle played on a square, two-dimensional board with numbered tiles that slide. The goal of this puzzle is to arrange the board's tiles from smallest to largest, left to right, top to bottom, with an empty space in board's bottom-right corner, as in the below.



Sliding any tile that borders the board's empty space in that space constitutes a "move." Although the configuration above depicts a game already won, notice how the tile numbered 12 or the tile numbered 15 could be slid into the empty space. Tiles may not be moved diagonally, though, or forcibly removed from the board.

Although other configurations are possible, we shall assume that this game begins with the board's tiles in reverse order, from largest to smallest, left to right, top to bottom, with an empty space in the board's bottom-right corner. **If, however, and only if the board contains an odd number of tiles (i.e., the height and width of the board are even), the positions of tiles numbered 1 and 2 must be swapped, as in the below.** The puzzle is solvable from this configuration.



Okay, navigate your way to `~/workspace/pset3/fifteen`, and take a look at `fifteen.c`. Within this file is an entire framework for the Game of Fifteen. The challenge up next is to complete this game's implementation.

But first go ahead and compile the framework. (Can you figure out how?) And, even though it's not yet finished, go ahead and run the game. (Can you figure out how?) Odds are you'll want to run it in a larger terminal window than usual, which you can open clicking the green plus (+) next to one of your code tabs and clicking **New Terminal**. Alternatively, you can full-screen the terminal window toward the bottom of CS50 IDE's UI (within the UI's "console") by clicking the **Maximize** icon in the console's top-right corner.

Anyhow, it appears that the game is at least partly functional. Granted, it's not much of a game yet. But that's where you come in!

questions

Read over the code and comments in `fifteen.c` and then answer the questions below in `questions.txt`, which is a (nearly empty) text file that we included for you inside of the distro's `fifteen` directory. No worries if you're not quite sure how `fprintf` or `fflush` work; we're simply using those to automate some testing.

1. Besides 4×4 (which are Game of Fifteen's dimensions), what other dimensions does the framework allow?
2. With what sort of data structure is the game's board represented?
3. What function is called to greet the player at game's start?
4. What functions do you apparently need to implement?

fifteen

Alright, get to it, implement this game. Remember, take "baby steps." Don't try to bite off the entire game at once. Instead, implement one function at a time and be sure that it works before forging ahead. In particular, we suggest that you implement the framework's functions in this order: `init`, `draw`, `move`, `won`. Any design decisions not explicitly prescribed herein (e.g., how much space you should leave between numbers when printing the board) are intentionally left to you. Presumably the board, when printed, should look something like the below, but we leave it to you to implement your own vision.

```
15 14 13 12
11 10  9  8
 7  6  5  4
 3  1  2  _
```

Incidentally, recall that the positions of tiles numbered 1 and 2 should only start off swapped (as they are in the 4×4 example above) if the board has an odd number of tiles (as does the 4×4 example above). If the board has an even number of tiles, those positions should not start off swapped. And so they do not in the 3×3 example below:

```
8  7  6
5  4  3
2  1  _
```

Here, now, is Zamyla:

fifteen



To test your implementation of `fifteen`, you can certainly try playing it. (Know that you can force your program to quit by hitting ctrl-c.) Be sure that you (and we) cannot crash your program, as by providing bogus tile numbers. And know that, much like you automated input into `find`, so can you automate execution of this game. In fact, in `~cs50/pset3` are `3x3.txt` and `4x4.txt`,

winning sequences of moves for a 3×3 board and a 4×4 board, respectively. To test your program with, say, the first of those inputs, execute the below.

```
./fifteen 3 < ~cs50/pset3/3x3.txt
```

Feel free to tweak the appropriate argument to `usleep` to speed up animation. In fact, you're welcome to alter the aesthetics of the game. For (optional) fun with "ANSI escape sequences," including color, take a look at our implementation of `clear` and check out http://isthe.com/chongo/tech/comp/ansi_escapes.html (http://isthe.com/chongo/tech/comp/ansi_escapes.html) for more tricks.

You're welcome to write your own functions and even change the prototypes of functions we wrote. But we ask that you not alter the flow of logic in `main` itself so that we can automate some tests of your program once submitted. In particular, `main` must only return `0` if and when the user has actually won the game; non-zero values should be returned in any cases of error, as implied by our distribution code. If in doubt as to whether some design decision of yours might run counter to the staff's wishes, simply contact your teaching fellow.

If you'd like to play with the staff's own implementation of `fifteen`, you may execute the below.

```
~cs50/pset3/fifteen
```

If you'd like to see an even fancier version, one so good that it can play itself, try out our solution to the Hacker Edition by executing the below.

```
~cs50/hacker3/fifteen
```

Instead of typing a number at the game's prompt, type `GOD` instead. Neat, eh?

And if you'd like to check the correctness of your program officially with `check50`, you may execute the below. **Note that `check50` assumes that your board's blank space is implemented in `board` as `0`; if you've chosen some other value, best to change to `0` for `check50`'s sake. Also note that `check50` assumes that you're indexing into `board` a la `board[row][column]`, not `board[column][row]`.**

```
check50 2015.fall.pset3.fifteen fifteen.c
```

How to Submit

Step 1 of 2

1. When ready to submit, log into [CS50 IDE \(https://cs50.io/\)](https://cs50.io/).
2. Toward CS50 IDE's top-left corner, within its "file browser" (not within a terminal window), control-click or right-click your `pset3` folder and then select **Download**. You should find that your browser has downloaded `pset3.tar.gz`, a "gzipped tarball" that's similar in spirit to a ZIP file.
3. In a separate tab or window, log into [CS50 Submit \(http://cs50.edx.org/submit\)](http://cs50.edx.org/submit), logging in if prompted.
4. Click **Submit** toward the window's top-left corner.
5. Under **Problem Set 3** on the screen that appears, click **Upload New Submission**.
6. On the screen that appears, click **Add files...**. A window entitled **Open Files** should appear.

7. Navigate your way to `pset3.tar.gz`. Odds are it's in your **Downloads** folder or wherever your browser downloads files by default. Once you find `pset3.tar.gz`, click it once to select it, then click **Open** (or the like).
8. Click **Start upload** to upload all of your files at once to CS50's servers.
9. On the screen that appears, you should see a window with **No File Selected**. If you move your mouse toward the window's lefthand side, you should see a list of the files you uploaded. Click each to confirm the contents of each. (No need to click any other buttons or icons.) If confident that you submitted the files you intended, consider your source code submitted! If you'd like to re-submit different (or modified) files, simply return to [CS50 Submit \(http://cs50.edx.org/submit\)](http://cs50.edx.org/submit) and repeat these steps. You may re-submit as many times as you'd like; we'll grade your most recent submission, so long as it's before the deadline.

Step 2 of 2

Head to <http://cs50.edx.org/2016/psets/3/> (<http://cs50.edx.org/2016/psets/3/>) where a short form awaits. Once you have submitted that form (as well as your source code), you are done! If you end up resubmitting your files (per step 1 of 1), no need to resubmit the form.

This was Problem Set 3.